



LEARNING

Application Architecture

Ruby

Java

Node.js

Python

PHP

Clojure

Scala

Go

Mobile

Database

Getting Started with Django on Heroku

🕒 Last updated 09 September 2015

☰ Table of Contents

- Prerequisites
- Start a Django app inside a Virtualenv
- Declare process types with Procfile
- Specify dependencies with Pip
- Django settings
- Store your app in Git
- Deploy to Heroku
- Dyno sleeping and scaling
- Next steps

This quickstart will get you going with a Python/Django application that uses a Postgres

database, deployed to Heroku. For other Python apps, see [Getting Started with Python on Heroku](#). For general information on how to develop and architect apps for use on Heroku, see [Architecting Applications for Heroku](#).



If you have questions about Python on Heroku, consider discussing it in the [Python on Heroku forums](#). Both Heroku and community-based Python experts are available.

Prerequisites

- The Heroku Toolbelt, as described in [Getting Started with Python](#).
- Installed [Python](#) and [Virtualenv](#) in a unix-style environment. See [this guide](#) for guidance.
- An installed version of [Postgres](#) to test locally.
- A Heroku user account. [Signup is free and instant](#).

Start a Django app inside a Virtualenv

First, we'll create an empty top-level directory for our project:

```
$ mkdir hellodjango && cd hellodjango
```



Make sure you're using the latest virtualenv release. If you're using a version that comes with Ubuntu, you may need to add the `--no-site-packages` flag.

Next, we'll create a Python [Virtualenv](#) (v1.0+):

```
$ virtualenv venv  
New python executable in venv/bin/python  
Installing setuptools, pip...done.
```

To use the new virtualenv, we need to activate it. (You must source the virtualenv environment for each terminal session where you wish to run your app.)

```
$ source venv/bin/activate
```

Next, install our application's dependencies with [pip](#). In this case, we will be installing **django-toolbelt**, which includes all of the packages we need:

- Django (the web framework)
- Gunicorn (WSGI server)
- dj-database-url (a Django configuration helper)

- dj-static (a Django static file server)

From your virtualenv:

```
$ pip install django-toolbelt
Installing collected packages: Django, psycopg2, gunicorn, dj-database-url, dj-static
...
Successfully installed Django psycopg2 gunicorn dj-database-url dj-static static
Cleaning up...
```

Now that we have a clean Python environment to work in, we'll create our simple Django application.



Don't forget the `.` at the end. This tells Django to put the extract the into the current directory, instead of putting it in a new subdirectory.

```
$ django-admin.py startproject hellodjango .
```

Declare process types with Procfile

Use a [Procfile](#), a text file in the root directory of your application, to explicitly declare what

command should be executed to start a web [dyno](#). In this case, you need to execute Gunicorn with a few arguments.

Here's a `Procfile` for our new app. It should be called `Procfile` and live at the root directory of our project:

Procfile

```
web: gunicorn hellodjango.wsgi --log-file -
```

You can now start the processes in your Procfile locally using the `heroku local` command, installed as part of the Heroku Toolbelt:

```
$ heroku local
2013-04-03 16:11:22 [8469] [INFO] Starting gunicorn 0.17.2
2013-04-03 16:11:22 [8469] [INFO] Listening at: http://127.0.0.1:8000 (8469)
```

Make sure things are working properly `curl` or a web browser, then Ctrl-C to exit.

Specify dependencies with Pip

Heroku recognizes Python applications by the existence of a `requirements.txt` file in the root of a repository. This simple format is used by most Python projects to specify the external

Python modules the application requires.

Pip has a nice command (`pip freeze`) that will generate this file for us:

```
$ pip freeze > requirements.txt
```

requirements.txt

```
Django==1.7  
dj-database-url==0.3.0  
dj-static==0.0.6  
gunicorn==19.1.1  
psycopg2==2.5.1  
static==0.4  
wsgiref==0.1.2
```

Pip can also be used for advanced dependency management. See [Python Dependencies via Pip](#) to learn more.

Django settings

Next, configure the application for the Heroku environment, including [Heroku's Postgres](#) database. The `dj-database-url` module will parse the value of the `DATABASE_URL` [environment variable](#) and convert them to something Django can understand.

Make sure 'dj-database-url' is in your requirements file, then add the following to the bottom of your `settings.py` file:

settings.py

```
# Parse database configuration from $DATABASE_URL
import dj_database_url
DATABASES['default'] = dj_database_url.config()

# Honor the 'X-Forwarded-Proto' header for request.is_secure()
SECURE_PROXY_SSL_HEADER = ('HTTP_X_FORWARDED_PROTO', 'https')

# Allow all host headers
ALLOWED_HOSTS = ['*']

# Static asset configuration
import os
BASE_DIR = os.path.dirname(os.path.abspath(__file__))
STATIC_ROOT = 'staticfiles'
STATIC_URL = '/static/'

STATICFILES_DIRS = (
    os.path.join(BASE_DIR, 'static'),
)
```

With these settings available, you can add the following code to `wsgi.py` to serve static files in production:

wsgi.py

```
from django.core.wsgi import get_wsgi_application
from dj_static import Cling

application = Cling(get_wsgi_application())
```

Store your app in Git

Now that we've written and tested our application, we need to store the project in a [Git](#) repository.

Since our current directory contains a lot of extra files, we'll want to configure our repository to ignore these files with a `.gitignore` file:



GitHub provides an excellent [Python gitignore file](#) that can be [installed system-wide](#).

.gitignore

```
venv
*.pyc
staticfiles
```

Next, we'll create a new git repository and save our changes.

```
$ git init
Initialized empty Git repository in /Users/kreitz/hellodjango/.git/
$ git add .
$ git commit -m "my django app"
[master (root-commit) 2943412] my django app
7 files changed, 230 insertions(+)
create mode 100644 .gitignore
create mode 100644 Procfile
create mode 100644 hellodjango/__init__.py
create mode 100644 hellodjango/settings.py
create mode 100644 hellodjango/urls.py
create mode 100644 hellodjango/wsgi.py
create mode 100644 manage.py
create mode 100644 requirements.txt
```

Deploy to Heroku

The next step is to push the application's repository to Heroku. First, we have to get a place to push to from Heroku. We can do this with the `heroku create` command:

```
$ heroku create
Creating simple-spring-9999... done, stack is cedar-14
http://simple-spring-9999.herokuapp.com/ | git@heroku.com:simple-spring-9999.git
Git remote heroku added
```

This automatically added the Heroku remote for our app (`git@heroku.com:simple-spring-9999.git`) to our repository. Now we can do a simple `git push` to deploy our application:

```
$ git push heroku master
Counting objects: 11, done.
Delta compression using up to 4 threads.
Compressing objects: 100% (9/9), done.
Writing objects: 100% (11/11), 4.01 KiB, done.
Total 11 (delta 0), reused 0 (delta 0)
-----> Python app detected
-----> No runtime.txt provided; assuming python-2.7.6.
-----> Preparing Python runtime (python-2.7.6)
-----> Installing Distribute (0.6.36)
-----> Installing Pip (1.3.1)
-----> Installing dependencies using Pip (1.3.1)
      Downloading/unpacking Django==1.5 (from -r requirements.txt (line 1))
      ...
      Successfully installed Django psycopg2 gunicorn dj-database-url dj-static stat
      Cleaning up...
-----> Collecting static files
      0 static files copied.

-----> Discovering process types
      Procfile declares types -> web

-----> Compiled slug size is 29.5MB
-----> Launching... done, v6
      http://simple-spring-9999.herokuapp.com deployed to Heroku

To git@heroku.com:simple-spring-9999.git
```

```
* [new branch]      master -> master
```

Visit your application

You've deployed your code to Heroku, and specified the process types in a `Procfile`. You can now instruct Heroku to execute a process type. Heroku does this by running the associated command in a `dyno` - a lightweight container which is the basic unit of composition on Heroku.

Let's ensure we have one dyno running the `web` process type:

```
$ heroku ps:scale web=1
```

You can check the state of the app's dynos. The `heroku ps` command lists the running dynos of your application:

```
$ heroku ps
=== web: `gunicorn hellodjango.wsgi`
web.1: up for 10s
```

Here, one dyno is running.

We can now visit the app in our browser with `heroku open`.

```
$ heroku open  
Opening simple-spring-9999.herokuapp.com... done
```

You should see the satisfying “It worked!” Django welcome page.

Dyno sleeping and scaling

By default, your app is deployed on a free dyno. Free dynos will sleep after a half hour of inactivity and they can be active (receiving traffic) for no more than 18 hours a day before [going to sleep](#). If a free dyno is sleeping, and it hasn’t exceeded the 18 hours, any web request will wake it. This causes a delay of a few seconds for the first request upon waking. Subsequent requests will perform normally.

To avoid dyno sleeping, you can upgrade to a hobby or professional dyno type as described in the [Dyno Types](#) article. For example, if you migrate your app to a professional dyno, you can easily scale it by running a command telling Heroku to execute a specific number of dynos, each running your web process type.

View the logs

Heroku treats logs as streams of time-stamped events aggregated from the output streams

of all the dynos running the components of your application. Heroku's [Logplex](#) provides a single channel for all of these events.

View information about your running app using one of the [logging commands](#), `heroku logs`:

```
$ heroku logs
2012-04-06T19:38:25+00:00 heroku[web.1]: State changed from created to starting
2012-04-06T19:38:29+00:00 heroku[web.1]: Starting process with command `unicorn hell
2012-04-06T19:38:29+00:00 app[web.1]: Validating models...
2012-04-06T19:38:29+00:00 app[web.1]:
2012-04-06T19:38:29+00:00 app[web.1]: 0 errors found
2012-04-06T19:38:29+00:00 app[web.1]: Django version 1.5, using settings 'hellodjango
2012-04-06T19:38:29+00:00 app[web.1]: Development server is running at http://0.0.0.0
2012-04-06T19:38:29+00:00 app[web.1]: Quit the server with CONTROL-C.
2012-04-06T19:38:30+00:00 heroku[web.1]: State changed from starting to up
2012-04-06T19:38:32+00:00 heroku[slugc]: Slug compilation finished
```

Syncing the database

The `heroku run` command lets you run [one-off dynos](#). You can use this to sync the Django models with the database schema:

```
$ heroku run python manage.py syncdb
Running python manage.py syncdb attached to terminal... up, run.1
Creating tables ...
```

```
Creating table auth_permission
Creating table auth_group_permissions
Creating table auth_group
Creating table auth_user_groups
Creating table auth_user_user_permissions
Creating table auth_user
Creating table django_content_type
Creating table django_session
Creating table django_site
```

```
You just installed Django's auth system, which means you don't have any superusers defined.
Would you like to create one now? (yes/no): yes
Username (leave blank to use 'u53976'): kenneth
Email address: kenneth@heroku.com
Password:
Password (again):
Superuser created successfully.
Installing custom SQL ...
Installing indexes ...
Installed 0 object(s) from 0 fixture(s)
```

Using the Django shell

Similarly, you can use `heroku run` to get a Django shell for executing arbitrary code against your deployed app:

```
$ heroku run python manage.py shell
```

```
Running python manage.py shell attached to terminal... up, run.1
Python 2.7.6 (default, Jan 16 2014, 02:39:37)
[GCC 4.4.3] on linux2
Type "help", "copyright", "credits" or "license" for more information.
(InteractiveConsole)
>>> from django.contrib.auth.models import User
>>> User.objects.all()
[<User: kenneth>]
```

Next steps

Now that you've deployed your first Django application to Heroku, it's time to take the next step! If you'd like to learn more about Heroku, these articles are a great place to start.

Heroku Reference

- [How Heroku Works](#)
- [Heroku Reference Documentation](#)

Python Reference

- [Django and Static Assets](#)
- [Using Celery on Heroku](#)

- [Deploying Python Applications with Gunicorn](#)
- [Specifying a Python Runtime](#)
- [Python Dependencies via Pip](#)
- [Background Tasks in Python with RQ](#)



django

python

Keep reading



[Python](#)



[Dynos and the Dyno Manager](#)



[Getting Started with Python on Heroku](#)



[Python Dependencies via Pip](#)



Feedback

[Log in to submit feedback.](#)

[< Direct to S3 File Uploads in Python](#)

[Getting Started with Python on Heroku >](#)

heroku.com

[Privacy Policy](#)